Machine Learning in Complex Domains: Assignment 2

Daniel Deutsch, Dan Crankshaw, Ryan Cotterell

3 Problem Set

3.1 Learning: Parameter Estimation

3.1.5 Analytical Questions

1. Overfitting typically occurs in a complex model when the number of parameters is much larger than the number of observations. When you overfit the data, the model is more likely to exaggerate noise in the data instead of lessening its effect. In this problem, consider the case if the robot reaches location (i, j) at time t. If the robot does not observe a wall to the north when it should, the model will learn that it should never observe a wall to the north when at location (i, j) at time t. We can think of this mistake the model made as noise. The model learned the noise to be true, so whenever the robot is at (i, j) at time t, it will always predict the robot will observe no wall. This noise in the data has been exaggerated to always be true.

   Therefore, when you are dealing with many parameters that are independent with respect to a specific variable, it is important to use parameter sharing so this situation does not happen and you lessen the effect of noise in the data. In this assignment, we say that a robot observing a wall is independent of the time it happens so that the model does not learn parameters for each time step and learn the noise into the model.

2. Suppose that in our model, the floor at location (i, j) is very sticky. When the robot tries to leave this location, it is more likely not to move compared to other locations on he map. When we use parameter sharing, we lose the information about moving in direction with respect to each position, and we can no longer encode that information. Parameter sharing makes the assumption that the probability of moving in a specific direction is the same for every position, and if that assumption is not true, parameter sharing cannot learn this model.

3. In order to combine these two approaches, you need to be able to detect an anomaly within the data, like a robot not being able to move as well in a specific location compared to the rest. A way to do this is to compute the statistics for every location individually. In our problem, this would be counting the number of times the robot is able to move in each direction in every location. Say that the sticky location is at (i, j).

   To detect that this is an anomaly, we will iterate over all of the positions, leaving 1 position out each time. Compute the average success of leaving the location and the rest of the locations. If the probability of leaving the sticky square is significantly different than from the rest of the locations, then leave that location out of the parameter sharing because it is an anomaly. Then you can learn parameters for each of the specific anomalies and shared parameters for the rest.

   With this approach, you can learn the specific parameters for each space that you need to and general parameters for when it does not matter. In this way, you have taken the advantages of both solutions.

3.2 Inference

3.2.1 Analytical Questions: Clique Tree

1. The process we used to create the clique tree is drawn out graphically in the file `CliqueTree Building.pdf`. The strategy we took was to try and do a small example (when the number of steps = 3), then try to generalize that clique tree to any arbitrary number of time steps.

   We made it as the instructions said: we converted the original graph into an undirected graph by marrying the parents of nodes at a v-structure; we converted the graph to a chordal graph, and verified that it was a valid chordal graph; we extracted the maximal cliques from the graph and formed the cluster graph over these cliques; finally, we found a maximal spanning tree that we thought made sense intuitively and could generalize. We tried to select the MST such that there was a pattern in the structure that we could generalize.

   When we were selecting the MST, we could have selected a different, but equally valid MST that would have produced another clique tree that is different from what we used. The ordering which we selected nodes in the algorithm to create a chordal graph was also important because if you did not select the nodes in an intelligent order, it was possible that the chordal graph would have been connected in a strange way. For example, there was a possibility that `PositionRow_t-1` would have been directly connected to `PositionRow_t+1`. If we allowed for that to happen, I believe that our resulting clique tree might be larger, more complicated, or less intuitive.

2. We have verified that the running intersection property holds. We wrote a program in `test.RunningIntersectionChecker.java` that takes a list of the the variables and finds all of the cliques whose scope contains them. Then we ran a breadth-first search to find a path, only adding an edge to the queue if the other clique contained that same variable. If for all variables and for all of the cliques contain them in their scope, there exists a path from each clique to every other clique such that all of the cliques along that path contain the variable in their scope, then the running intersection property holds. Our program verified that our clique trees are correct.

3. We made the `cliquetree` file by generalizing the clique tree that we came up with in `CliqueTree Building.pdf`. We then wrote a Python script to output a clique tree for an arbitrary number of time steps and landmarks, the only variables that change between each clique tree. The script that we wrote is called `create-clique.py`.

3.2.5 Empirical Questions: Message Passing

1. In order to find the distribution over the final position of the robot, we added all of the evidence for the respective files to the model after calibration. Then we ran the queries

   ```
   PositionRow_t=1,PositionCol_t=1
   PositionRow_t=1,PositionCol_t=2
   ...
   PositionRow_t=10,PositionCol_t=10
   ```

   for $t = 9, 99$ and 999. These query files are located in the subdirectory `queries` and are `problem1-t10.txt`, `problem1-t100.txt`, and `problem1-t1000.txt`.

   For the `network-grid10x10-t10.txt` file, our resulting probabilities were:

10	0.00	0.01	0.48	1.57	0.89	1.14	0.16	2.77	1.03	0.32
9	0.02	0.01	0.50	1.55	0.96	10.27	1.97	0.27	2.29	0.13
8	0.02	0.40	0.62	8.51	3.09	0.14	0.17	0.24	1.21	1.82
7	0.32	0.28	0.32	0.19	0.20	1.81	0.26	0.61	0.16	0.09
6	0.01	0.04	0.01	0.19	0.28	1.08	0.10	0.05	0.02	0.43
5	0.11	2.50	0.36	0.04	0.59	1.88	0.23	0.02	0.10	0.31
4	0.77	2.70	0.04	0.34	2.71	1.57	0.11	0.12	0.05	0.17
3	0.04	0.03	0.06	0.18	0.63	0.34	2.10	0.01	0.68	5.55
2	0.08	0.03	0.02	0.17	0.25	3.56	0.74	2.04	2.04	2.01
1	0.02	0.00	0.41	1.78	4.26	0.54	3.39	0.33	4.96	0.10
	1	2	3	4	5	6	7	8	9	10

Where the value at row i and column j is the percent probability of the robot ending up at position (i, j).


Figure 1: The contour plot for t10

As you can see from the contour plot for t10 in Figure 1, the most likely location for the robot at time $t = 9$ is the top middle of the map.

For the `network-grid10x10-t100.txt` file, our resulting probabilities were:

10	0.01	0.02	0.67	3.48	2.31	2.99	0.35	2.94	2.33	0.35
9	0.12	0.01	0.63	1.96	0.49	0.59	2.19	0.31	2.18	0.40
8	0.05	0.61	0.65	0.45	2.68	0.51	0.39	0.61	3.03	3.50
7	0.35	0.03	0.44	0.28	0.36	2.69	0.39	0.57	0.36	0.16
6	0.02	0.06	0.01	0.39	0.81	3.11	0.18	0.10	0.07	0.42
5	0.10	0.09	0.39	0.05	2.59	2.11	0.47	0.07	0.43	0.84
4	2.40	0.10	0.03	0.40	3.72	3.66	0.70	0.42	0.07	0.47
3	0.06	0.05	0.12	0.36	0.55	0.81	3.03	0.02	3.81	0.26
2	0.12	0.06	0.15	0.42	0.61	2.56	0.06	3.04	0.10	3.91
1	0.07	0.01	0.50	3.10	2.12	0.45	3.10	0.18	3.32	0.31
	1	2	3	4	5	6	7	8	9	10

and the corresponding contour plot is in Figure 2. From the Figure, you can see that the robot is likely in the middle or the bottom right of the map.


Figure 2: The contour plot for t100

We were unable to have our `bayes-query-sp` program terminate on the largest file, `network-grid10x10-t1000.txt`.

2. The set of queries that can be located in file `problem2.txt` ask for the probability of:

```
ObserveLandmark1_N_8 evidence
ObserveLandmark1_E_8
ObserveLandmark1_S_8
ObserveLandmark1_W_8
```

Those probabilities were equal to 0.1703, 0.1707 0.1700, and 01698, respectively.

3. In order to compute this probability, we know that we can't simply perform the query because there is no clique in the clique tree with all four position vectors. In order to get around that, we're going to use the chain rule of probability:

$$P(\text{Position\_0, Position\_1} \,|\, \text{evidence}) = P(\text{Poisition\_1} \,|\, \text{Position\_0})P(\text{Position\_0} \,|\, \text{evidence})$$

This query can be found in the file `problem3.txt`.

The resulting probabilities of the two queries were 0.002507 and 0.001716. Thus, the combined result is 4.302E-6

4. Since we were instructed now to implement algorithm 10.4, which would have allowed to compute the joint probability of random variables assigned to different cliques, we resorted to the chain rule of probability.

$$p(A, B, \ldots, Z) = p(A)p(B|A) \ldots p(Z|A, B \ldots, Y)$$

Thus we could query each individual variable separately conditioning on the previous variables. We can then simple that the product ($\otimes$ of the semiring which is $+$ in logspace). After words we took the -log of the probability and reported it.

The query file that was written for this is located in the subdirectory `queries` with the file name `problem4.txt`. The result that we got after multiplying all of the subqueries together was 104.83 which is the log likelihood of the data.

5. In order to compare how much faster our program runs with the incremental updates, we took the queries from the previous problem, and altered them such that on every line, all of the evidence that we know to be true is given. We ran the queries from the previous problem (one piece of new evidence added per line) with the incremental updates enabled. This took 11.739 seconds as measured by `time`. Then we disabled the incremental updates in our program and ran the altered query file. This took 83.127 seconds as measured by `time`.

With the altered query file, we needed to add all of the evidence we knew at that line to be true to the right hand side of the query because the clique tree was reinitialized for every query execution. This was how we disabled the incremental updates. The results of both the enabled and disabled versions were the same, but the disabled version took much longer. The query file we ran with incremental updates enabled is called `problem5-enabled.txt`, and the file we ran with updates disabled is called `problem5-disabled.txt`. Both of these files can be found in the subdirectory `queries`.

Note, if you want to verify these results, you may have to contact us because in order to measure the time without incremental updates, we had to edit the source of our program.

6. There are many problems with this comparison for our particular case. Specifically, we implemented variable elimination instead of the brute force method for the first homework. Secondly, we did them in two different languages. We understand that the point of this assignment is to demonstrate to us the practical need for efficient algorithm en lieu of brute force. The experiment was designed as follows: we ran both Belief Propagation and Variable Elimination on `network-10-8.txt` and `grid10x10-t10.txt`. The BP ran instantly (on the order of milliseconds) on both maps whereas the Python implementation of Variable Elimination ran instantly (milliseconds) on the `network-10-8.txt` but ate up all the memory after 30 minutes on the `grid10x10-t10.txt`. In fact, it called the OOM-killer much to my dismay. What is more, the Python implementation ran in `grid10x10-t10.txt` ran in 10 milliseconds whereas the Java BP ran in 30 milliseconds on the same map; we attribute this to the start up time for JVM. I wish it were possible to do a more robust comparison, but I feel after talking to several of my classmates, I am thoroughly convinced that brute force is a bad approach.

3.2.6 Extra Credit: Max-Product Message Passing

1.

2.

3.3 Bayesian Score for Bayesian Networks

1.

$$P(\mathcal{D}|\mathcal{G}) = \prod_i \prod_{u_i \in Val Pa_{X_i}^{\mathcal{G}}} \frac{\Gamma(\alpha_{X_i|u_i}^{\mathcal{G}})}{\Gamma(\alpha_{X_i}^{\mathcal{G}} + M[u_i])} \prod_{x_i^j \in Val(X_i)} \frac{\Gamma(\alpha_{x_i^j} + M[x_i^j, u_i]}{\Gamma\alpha_{x_i^j|u_i}^{\mathcal{G}}}$$

$$\log(P(\mathcal{D}|\mathcal{G})) = \sum_i \sum_{u_i \in Val(Pa_{X_i}^{\mathcal{G}})} \log\left(\frac{\Gamma(\alpha_{X_i|u_i}^{\mathcal{G}})}{\Gamma(\alpha_{X_i}^{\mathcal{G}} + M[u_i])} \prod_{x_i^j \in Val(X_i)} \frac{\Gamma(\alpha_{x_i^j} + M[x_i^j, u_i]}{\Gamma\alpha_{x_i^j|u_i}^{\mathcal{G}}} \right)$$

$$\log(P(\mathcal{D}|\mathcal{G})) = \sum_i \sum_{u_i \in Val(Pa_{X_i}^{\mathcal{G}})} \left(\log\left(\frac{\Gamma(\alpha_{X_i|u_i}^{\mathcal{G}})}{\Gamma(\alpha_{X_i}^{\mathcal{G}} + M[u_i])} \right) + \sum_{x_i^j \in Val(X_i)} \log\left(\frac{\Gamma(\alpha_{x_i^j} + M[x_i^j, u_i]}{\Gamma\alpha_{x_i^j|u_i}^{\mathcal{G}}} \right) \right)$$

$$\sum_i \sum_{u_i \in Val(Pa_{X_i}^{\mathcal{G}})} \left(\log\left(\frac{\sqrt{2\pi}\alpha_{X_i}^{\mathcal{G}}\left(\alpha_{X_i}^{\mathcal{G}} - \frac{1}{2}\right) e^{-\alpha_{X_i}^{\mathcal{G}}}}{\sqrt{2\pi}\left(\alpha_{X_i}^{\mathcal{G}} + M[u_i]\right)^{\left(\alpha_{X_i}^{\mathcal{G}} + M[u_i] - \frac{1}{2}\right)} e^{-\alpha_{X_i}^{\mathcal{G}} + M[u_i]}} \right) + \sum_{x_i^j \in Val(X_i)} \log\left(\frac{\sqrt{2\pi}\left(\alpha_{x_i^j} + M[}{\sqrt{2\pi}} \right.$$

$$\sum_i \sum_{u_i \in Val(Pa_{X_i}^{\mathcal{G}})} \left(\log\left(\frac{\alpha_{X_i}^{\mathcal{G}}\left(\alpha_{X_i}^{\mathcal{G}} - \frac{1}{2}\right) e^{-M[u_i]}}{\left(\alpha_{X_i}^{\mathcal{G}} + M[u_i]\right)^{\left(\alpha_{X_i}^{\mathcal{G}} + M[u_i] - \frac{1}{2}\right)}} \right) + \sum_{x_i^j \in Val(X_i)} \log\left(\frac{\left(\alpha_{x_i^j} + M[x_i^j, u_i]\right)^{\left(\alpha_{x_i^j} + M[x_i^j, u}\right.}{\left(\alpha_{x_i^j|u_i}^{\mathcal{G}}\right)^{\left(\alpha_{x_i^j|u_i}^{\mathcal{G}} - \right.}} \right.$$

We whiteboarded the above to get it in the form below. It is not difficult to see how this would happen. When you push in the log the exponents are going to come down with the log. We didn't show this because would take up a lot more space and is not very interesting. Latex's multiline mode messed up the parentheses.

$$\sum_i \sum_{u_i} (\alpha^G_{X_i|u_i} - \frac{1}{2}) \ln(\alpha^G_{X_i|u_i}) - (\alpha^G_{X_i|u_i} - \frac{1}{2}) \ln(\alpha^G_{X_i|u_i} + M[u_i]) + M[u_i] +$$

$$\sum_j \left((\alpha^G_{x_i^j|u_i} - \frac{1}{2}) \ln(\alpha^G_{x_i^j|u_i} + M[x_i^j, u_i]) + M[x_i^j, u_i] - (\alpha_{x_i^j|u_i} - \frac{1}{2}) \ln(\alpha_{x_i^j|u_i}) \right)$$

$$\sum_i \sum_{u_i} (\alpha^G_{X_i|u_i} - \frac{1}{2}) \ln(\alpha^G_{X_i|u_i}) - (\alpha^G_{X_i|u_i} - \frac{1}{2} + M[u_i]) \ln(\alpha^G_{X_i|u_i} + M[u_i]) +$$

$$\sum_j \left((\alpha^G_{x_i^j|u_i} - \frac{1}{2} - M[x_i^j, u_i]) \ln(\alpha^G_{x_i^j|u_i} + M[x_i^j, u_i]) - (\alpha_{x_i^j|u_i} - \frac{1}{2}) \ln(\alpha_{x_i^j|u_i}) \right)$$

$$\sum_i \sum_{u_i} (\alpha^G_{X_i|u_i} - \frac{1}{2}) \ln(\alpha^G_{X_i|u_i}) - (\alpha^G_{X_i|u_i} - \frac{1}{2}) \ln(\alpha^G_{X_i|u_i} + M[u_i]) + (M[u_i] \ln(\alpha^G_{X_i|u_i} + M[u_i]) -$$

$$\sum_j \left((\alpha^G_{x_i^j|u_i} - \frac{1}{2}) \ln(\alpha^G_{x_i^j|u_i} + M[x_i^j, u_i]) + M[x_i^j, u_i] \ln(\alpha^G_{x_i^j|u_i} + M[x_i^j, u_i]) - (\alpha_{x_i^j|u_i} - \frac{1}{2}) \ln(\alpha_{x_i^j|u_i}) \right)$$

$$\sum_i \sum_{u_i} (\alpha^G_{X_i|u_i} - \frac{1}{2}) \ln(\frac{\alpha^G_{X_i|u_i}}{\alpha^G_{X_i|u_i} + M[u_i]}) + (M[u_i] \ln(\alpha^G_{X_i|u_i} + M[u_i]) -$$

$$\sum_j \left((\alpha^G_{x_i^j|u_i} - \frac{1}{2}) \ln(\frac{\alpha^G_{x_i^j|u_i} + M[x_i^j, u_i]}{\alpha_{x_i^j|u_i}}) + M[x_i^j, u_i] \ln(\alpha^G_{x_i^j|u_i} + M[x_i^j, u_i]) \right)$$

$$\sum_i \sum_{u_i} (\alpha^G_{X_i|u_i} - \frac{1}{2}) \ln(\frac{\alpha^G_{X_i|u_i}}{\alpha^G_{X_i|u_i} + M[u_i]}) -$$

$$\sum_j \left((\alpha^G_{x_i^j|u_i} - \frac{1}{2}) \ln(\frac{\alpha^G_{x_i^j|u_i} + M[x_i^j, u_i]}{\alpha_{x_i^j|u_i}}) + M[x_i^j, u_i] \ln(\alpha^G_{x_i^j|u_i} + M[x_i^j, u_i]) - (M[u_{ij}] \ln(\alpha^G_{X_i|u_i} + M[u_i]) \right)$$

$$\sum_i \sum_{u_i} (\alpha^G_{X_i|u_i} - \frac{1}{2}) \ln(\frac{\alpha^G_{X_i|u_i}}{\alpha^G_{X_i|u_i} + M[u_i]}) -$$

$$\sum_j \left((\alpha^G_{x_i^j|u_i} - \frac{1}{2}) \ln(\frac{\alpha^G_{x_i^j|u_i} + M[x_i^j, u_i]}{\alpha_{x_i^j|u_i}}) + M[x_i^j, u_i] \ln(\alpha^G_{x_i^j|u_i} + M[x_i^j, u_i]) - (M[x_i^j, u_j] \ln(\alpha^G_{X_i|u_i} + M[u_i]) \right)$$

$$\sum_i \sum_{u_i} (\alpha^G_{X_i|u_i} - \frac{1}{2}) \ln(\frac{\alpha^G_{X_i|u_i}}{\alpha^G_{X_i|u_i} + M[u_i]}) -$$

$$\sum_j \left((\alpha^G_{x_i^j|u_i} - \frac{1}{2}) \ln(\frac{\alpha^G_{x_i^j|u_i} + M[x_i^j, u_i]}{\alpha_{x_i^j|u_i}}) + M[x_i^j, u_i] \ln(\frac{\alpha^G_{x_i^j|u_i} + M[x_i^j, u_i]}{\alpha^G_{X_i|u_i} + M[u_i]}) \right)$$

$$\sum_i \sum_{u_i} (\alpha^G_{X_i|u_i} - \frac{1}{2}) \ln(\frac{\alpha^G_{X_i|u_i}}{\alpha^G_{X_i|u_i} + M[u_i]}) -$$

$$\sum_j \left((\alpha^G_{x_i^j|u_i} - \frac{1}{2}) \ln(\frac{\alpha^G_{x_i^j|u_i} + M[x_i^j, u_i]}{\alpha_{x_i^j|u_i}}) + M[x_i^j, u_i] \ln(\frac{\alpha^G_{x_i^j|u_i} + M[x_i^j, u_i]}{\alpha^G_{X_i|u_i} + M[u_i]}) \right)$$

Thus by rewriting and using the definition of the Likelihood score on Page 792, we get the following. Note we have removed the α since they are constant are dwarfed by M.

$$\sum_i \sum_{u_i} \sum_j M[x_i^j, u_i] \ln(\frac{\alpha^G_{x_i^j|u_i} + M[x_i^j, u_i]}{\alpha^G_{X_i|u_i} + M[u_i]}) = \ell(\hat{\theta}_{\mathcal{G}} : \mathcal{D})$$

With some rewriting we get,

$$\ell(\hat{\theta}^{\mathcal{G}}) + \sum_i \sum_{u_i} (\alpha^G_{X_i|u_i} - \frac{1}{2}) \ln(\frac{\alpha^G_{X_i|u_i}}{\alpha^G_{X_i|u_i} + M[u_i]}) - \sum_j \left((\alpha^G_{x_i^j|u_i} - \frac{1}{2}) \ln(\frac{\alpha^G_{x_i^j|u_i} + M[x_i^j, u_i]}{\alpha_{x_i^j|u_i}}) \right)$$

We can merge the two sums since $\sum_j$ iterations $|X_i|$ times

$$\ell(\hat{\theta}^{\mathcal{G}}) + \sum_i \sum_{u_i} \sum_j \frac{1}{|X_i|} \left((\alpha^G_{X_i|u_i} - \frac{1}{2}) \ln(\frac{\alpha^G_{X_i|u_i}}{\alpha^G_{X_i|u_i} + M[u_i]}) \right) - \left((\alpha^G_{x_i^j|u_i} - \frac{1}{2}) \ln(\frac{\alpha^G_{x_i^j|u_i} + M[x_i^j, u_i]}{\alpha_{x_i^j|u_i}}) \right)$$

I couldn't really get the right part of this into the correct form, but intuitively it makes sense why it would be $\log(M) \dim(G)$. We are iterating $\dim G$ since by enumerating all the i, u_i, j we are enumerating all the parameters. Then we just have to show that rest grows at a rate proportional to $\log(M)$. This seems like it should be true, but that's about as far as I got.